

Search Engine Architecture

11. Big Data Processing Part Two



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

Noted slides adapted from Lin et al.'s Big Data Infrastructure, UMD Spring 2015 with cosmetic changes.

Today's Agenda

- Making Hadoop more efficient
- Dataflow languages
- What's next?

Hadoop is slow...



Source: Wikipedia (Tortoise) via Lin et al. Big Data Infrastructure, UMD Spring 2015.

A Major Step Backwards?

- MapReduce is a step backward in database access:
 - Schemas are good
 - Separation of the schema from the application is good
 - High-level access languages are good
- MapReduce is poor implementation
 - Brute force and only brute force (no indexes, for example)
- MapReduce is not novel
- MapReduce is missing features
 - Bulk loader, indexing, updates, transactions...

Hadoop vs. Databases: Grep




Figure 4: Grep Task Results – 535MB/node Data Set




Figure 5: Grep Task Results – 1TB/cluster Data Set

```
SELECT * FROM Data WHERE field LIKE '%XYZ%';
```

Hadoop vs. Databases: Select




Figure 6: Selection Task Results

```
SELECT pageURL, pageRank  
FROM Rankings WHERE pageRank > X;
```

Hadoop vs. Databases: Aggregation




Figure 7: Aggregation Task Results (2.5 million Groups)




Figure 8: Aggregation Task Results (2,000 Groups)

```
SELECT sourceIP, SUM(adRevenue)
FROM UserVisits GROUP BY sourceIP;
```

Hadoop vs. Databases: Join




Figure 9: Join Task Results

facebook®

Jeff Hammerbacher, Information Platforms and the Rise of the Data Scientist.
In, *Beautiful Data*, O'Reilly, 2009.

"On the first day of logging the Facebook clickstream, more than 400 gigabytes of data was collected. The load, index, and aggregation processes for this data set really taxed the Oracle data warehouse. Even after significant tuning, we were unable to aggregate a day of clickstream data in less than 24 hours."



Why?

`Integer.parseInt`
`String.substring`

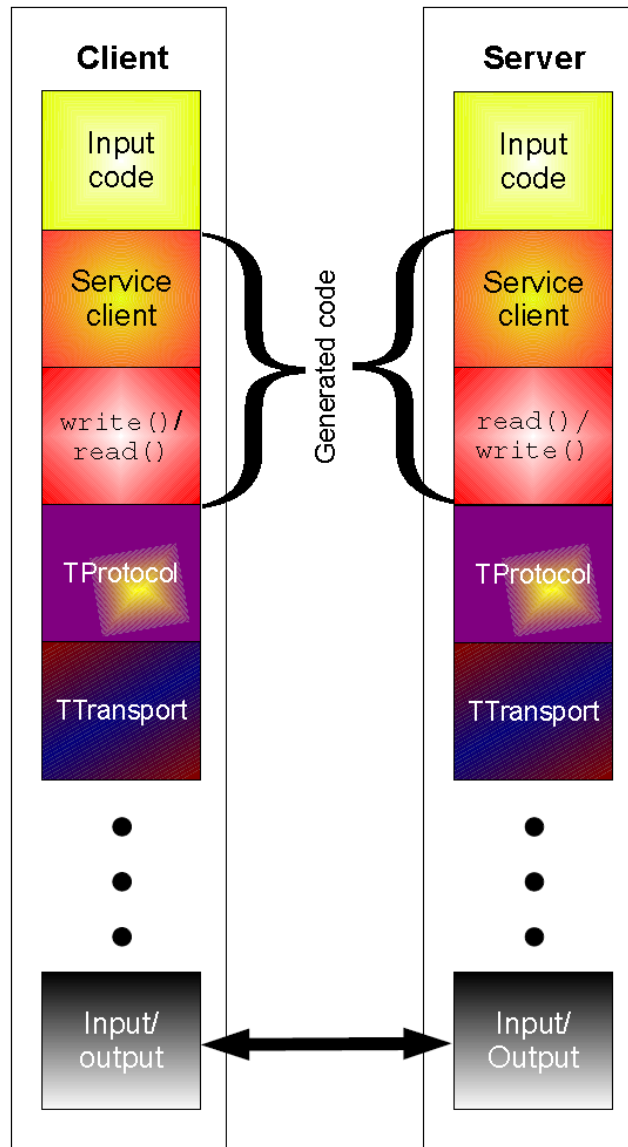
Schemas are a good idea!

- Parsing fields out of flat text files is slow
- Schemas define a contract, decoupling logical from physical

Thrift

- Originally developed by Facebook, now an Apache project
- Provides a DDL with numerous language bindings
 - Compact binary encoding of typed structs
 - Fields can be marked as optional or required
 - Compiler automatically generates code for manipulating messages
- Provides RPC mechanisms for service definitions
- Alternatives include protobufs, Avro, Parquet

Thrift



```
struct Tweet {  
  1: required i32 userId;  
  2: required string userName;  
  3: required string text;  
  4: optional Location loc;  
}  
  
struct Location {  
  1: required double latitude;  
  2: required double longitude;  
}
```

Dataflow Languages

Need for High-Level Languages

- Hadoop is great for large-data processing!
 - But writing Java programs for everything is verbose and slow
 - Data scientists don't want to write Java
- Solution: develop higher-level data processing languages
 - Hive: HQL is like SQL
 - Pig: Pig Latin is a bit like Perl

Hive and Pig

- Hive: data warehousing application in Hadoop
 - Query language is HQL, variant of SQL
 - Tables stored on HDFS with different encodings
 - Developed by Facebook, now open source
- Pig: large-scale data processing system
 - Scripts are written in Pig Latin, a dataflow language
 - Programmer focuses on data transformations
 - Developed by Yahoo!, now open source
- Common idea:
 - Provide higher-level language to facilitate large-data processing
 - Higher-level language “compiles down” to Hadoop jobs



Hive: Example


- Hive looks similar to an SQL database
- Relational join on two tables:
 - Table of word counts from Shakespeare collection
 - Table of word counts from the bible

```
SELECT s.word, s.freq, k.freq FROM shakespear s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```

the	25848	62394
I	23031	8854
and	19671	38985
to	18038	13526
of	16700	34654
a	14170	8057
you	12702	2720
my	11297	4135
in	10797	12445
is	8882	6884


Hive: Behind the Scenes

```
SELECT s.word, s.freq, k.freq FROM shakespeare s  
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1  
ORDER BY s.freq DESC LIMIT 10;
```



(Abstract Syntax Tree)

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF shakespeare s) (TOK_TABREF bible k) (= (. (TOK_TABLE_OR_COL s) word) (.  
(TOK_TABLE_OR_COL k) word)))) (TOK_INSERT (TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT (TOK_SELEXPR (.  
(TOK_TABLE_OR_COL s) word)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL s) freq)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL k) freq))) (TOK_WHERE  
(AND (>= (. (TOK_TABLE_OR_COL s) freq) 1) (>= (. (TOK_TABLE_OR_COL k) freq) 1))) (TOK_ORDERBY (TOK_TABSORTCOLNAMEDESC (.  
(TOK_TABLE_OR_COL s) freq))) (TOK_LIMIT 10)))
```



(one or more of MapReduce jobs)

Hive: Behind the Scenes

STAGE DEPENDENCIES:

Stage-1 is a root stage
Stage-2 depends on stages: Stage-1
Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1
Map Reduce

Alias -> Map Operator Tree:

```
s
  TableScan
  alias: s
  Filter Operator
  predicate:
    expr: (freq >= 1)
    type: boolean
  Reduce Output Operator
  key expressions:
    expr: word
    type: string
  sort order: +
  Map-reduce partition columns:
    expr: word
    type: string
  tag: 0
  value expressions:
    expr: freq
    type: int
    expr: word
    type: string
```

k

```
TableScan
alias: k
Filter Operator
predicate:
  expr: (freq >= 1)
  type: boolean
Reduce Output Operator
key expressions:
  expr: word
  type: string
sort order: +
Map-reduce partition columns:
  expr: word
  type: string
tag: 1
value expressions:
  expr: freq
  type: int
```

Reduce Operator Tree:

```
Join Operator
condition map:
  Inner Join 0 to 1
condition expressions:
  0 {VALUE._col0} {VALUE._col1}
  1 {VALUE._col0}
outputColumnNames: _col0, _col1, _col2
Filter Operator
predicate:
  expr: ((_col0 >= 1) and (_col2 >= 1))
  type: boolean
Select Operator
expressions:
  expr: _col1
  type: string
  expr: _col0
  type: int
  expr: _col2
  type: int
outputColumnNames: _col0, _col1, _col2
File Output Operator
compressed: false
GlobalTableId: 0
table:
  input format: org.apache.hadoop.mapred.SequenceFileInputFormat
  output format: org.apache.hadoop.hive.ql.io.HiveSequenceFileOutputFormat
```

Stage: Stage-2

Map Reduce

Alias -> Map Operator Tree:

hdfs://localhost:8022/tmp/hive-training/364214370/10002

Reduce Output Operator

key expressions:

expr: _col1
type: int

sort order: -
tag: -1

value expressions:

expr: _col0
type: string
expr: _col1
type: int
expr: _col2
type: int

Reduce Operator Tree:

Extract

Limit

File Output Operator

compressed: false
GlobalTableId: 0

table:

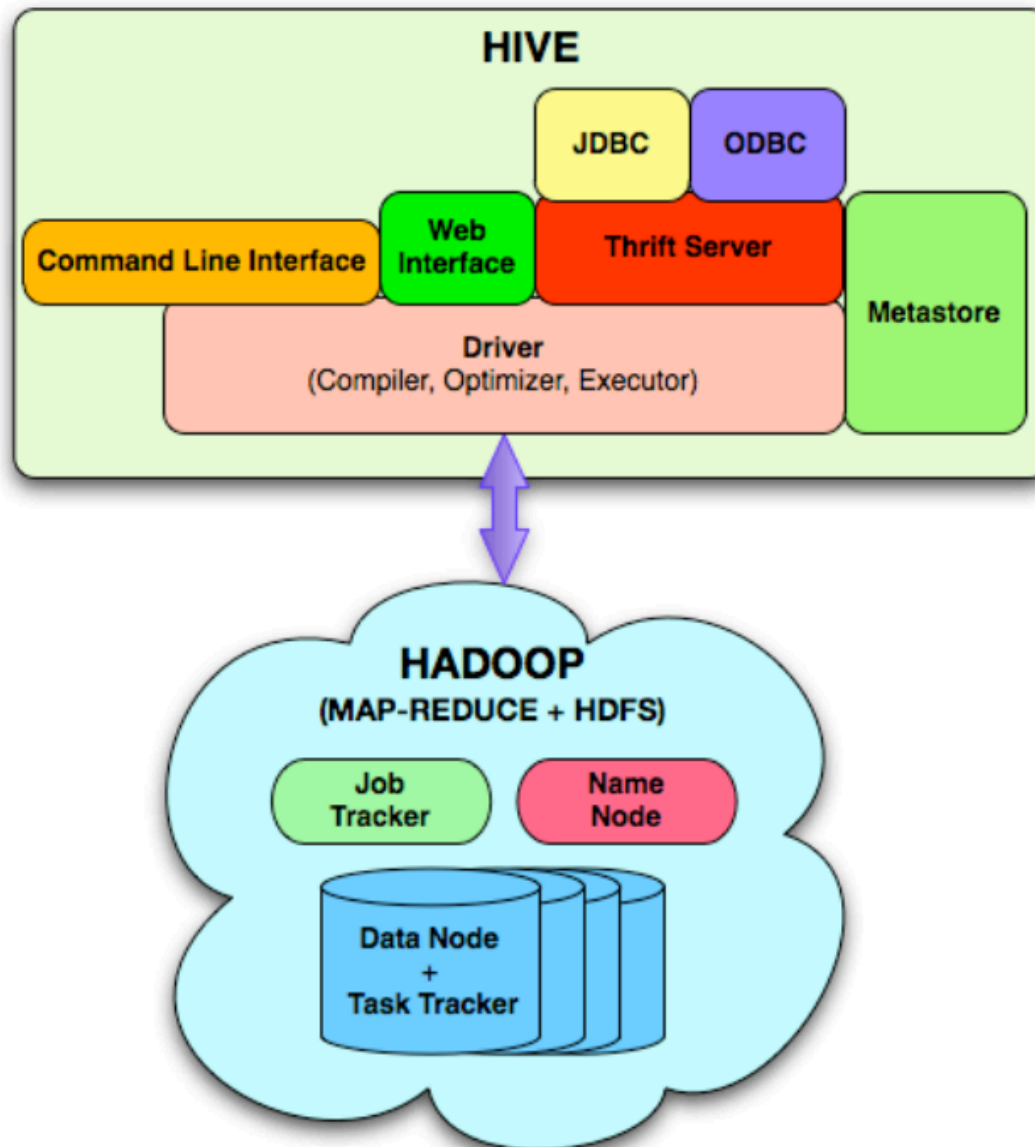
input format: org.apache.hadoop.mapred.TextInputFormat
output format: org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat

Stage: Stage-0

Fetch Operator

limit: 10

Hive Architecture



Hive Implementation

- Metastore holds metadata
 - Databases, tables
 - Schemas (field names, field types, etc.)
 - Permission information (roles and users)
- Hive data stored in HDFS
 - Tables in directories
 - Partitions of tables in sub-directories
 - Actual data in files

Pig!



Source: Wikipedia (Pig) via Lin et al. Big Data Infrastructure, UMD Spring 2015.

Pig: Example

Task: Find the top 10 most visited pages in each category

Visits

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	flickr.com	10:05
Fred	cnn.com	12:00



Url Info

Url	Category	PageRank
cnn.com	News	0.9
bbc.com	News	0.8
flickr.com	Photos	0.7
espn.com	Sports	0.9

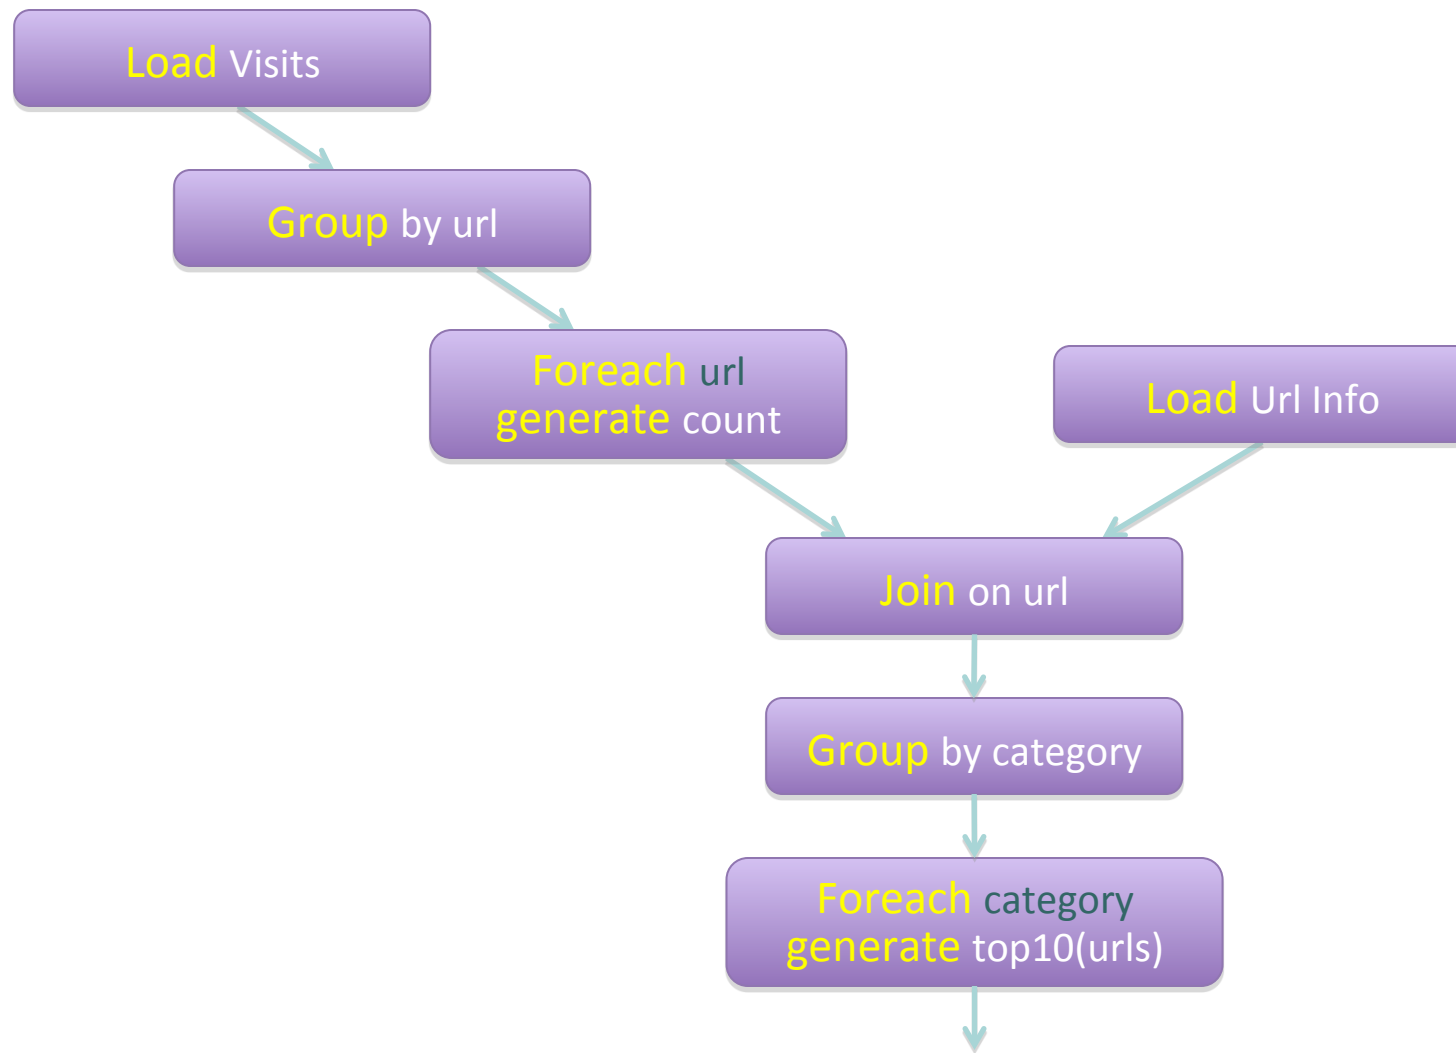


Pig Script

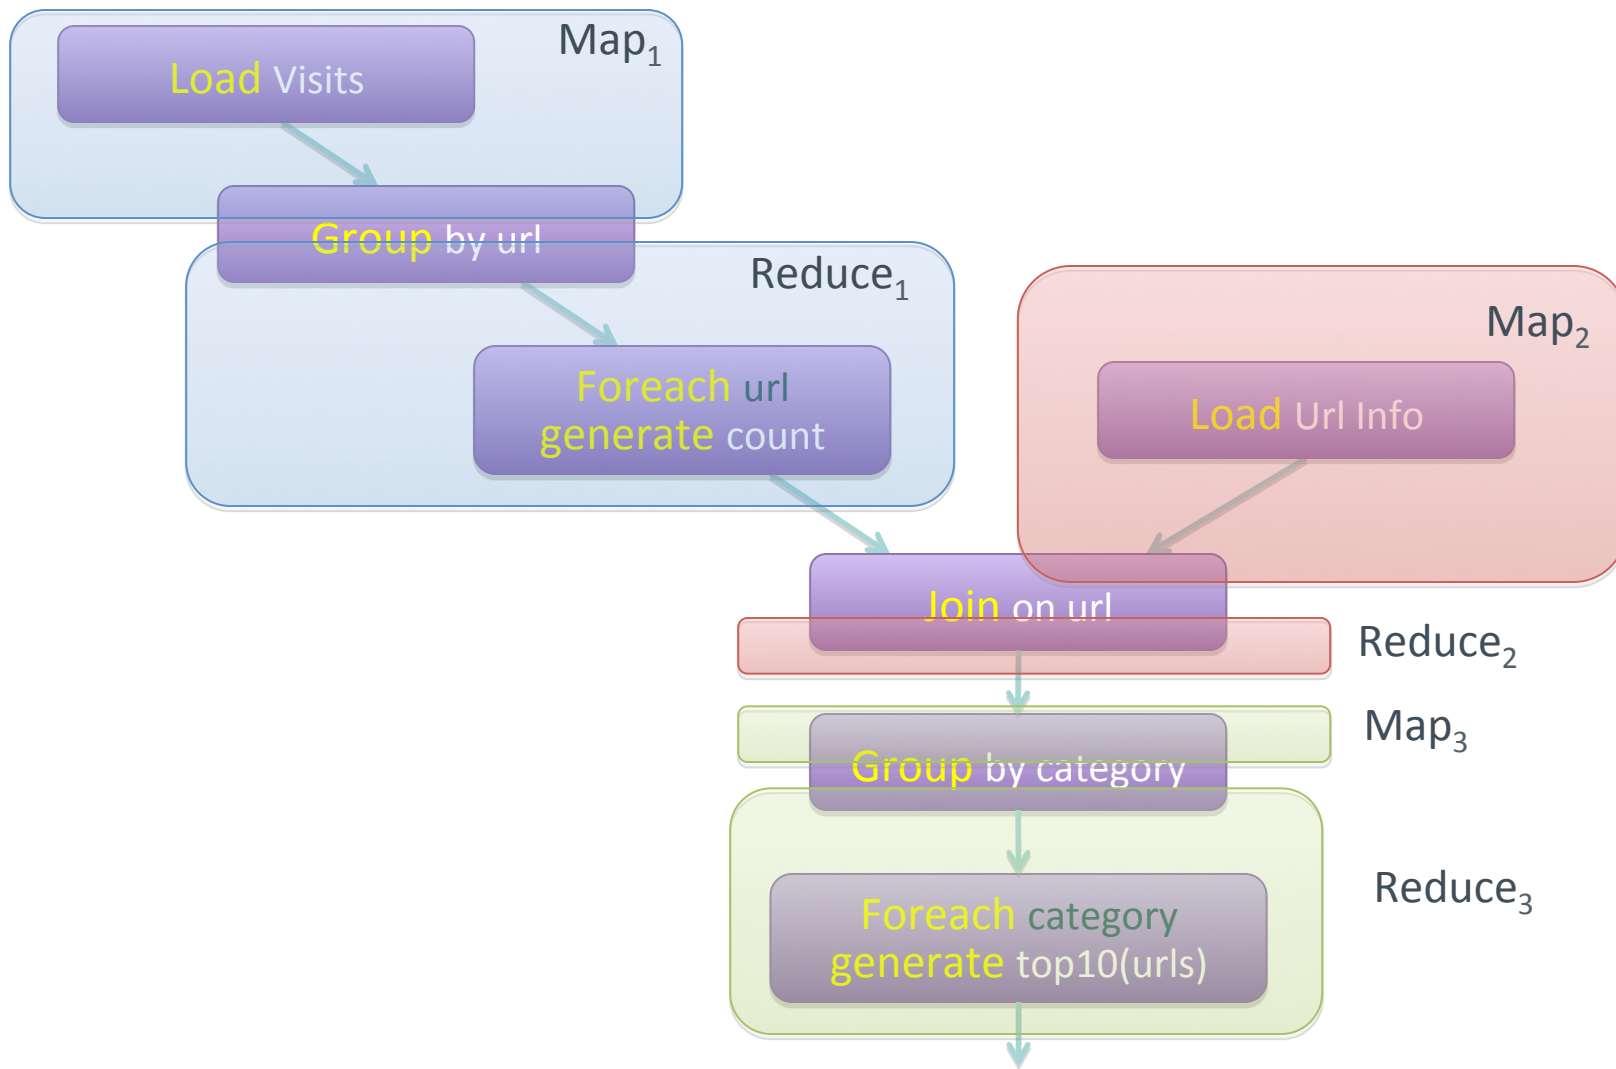
```
visits = load '/data/visits' as (user, url, time);
gVisits = group visits by url;
visitCounts = foreach gVisits generate url, count(visits);
urlInfo = load '/data/urlInfo' as (url, category, pRank);
visitCounts = join visitCounts by url, urlInfo by url;
gCategories = group visitCounts by category;
topUrls = foreach gCategories generate top(visitCounts,10);

store topUrls into '/data/topUrls';
```


Pig Query Plan

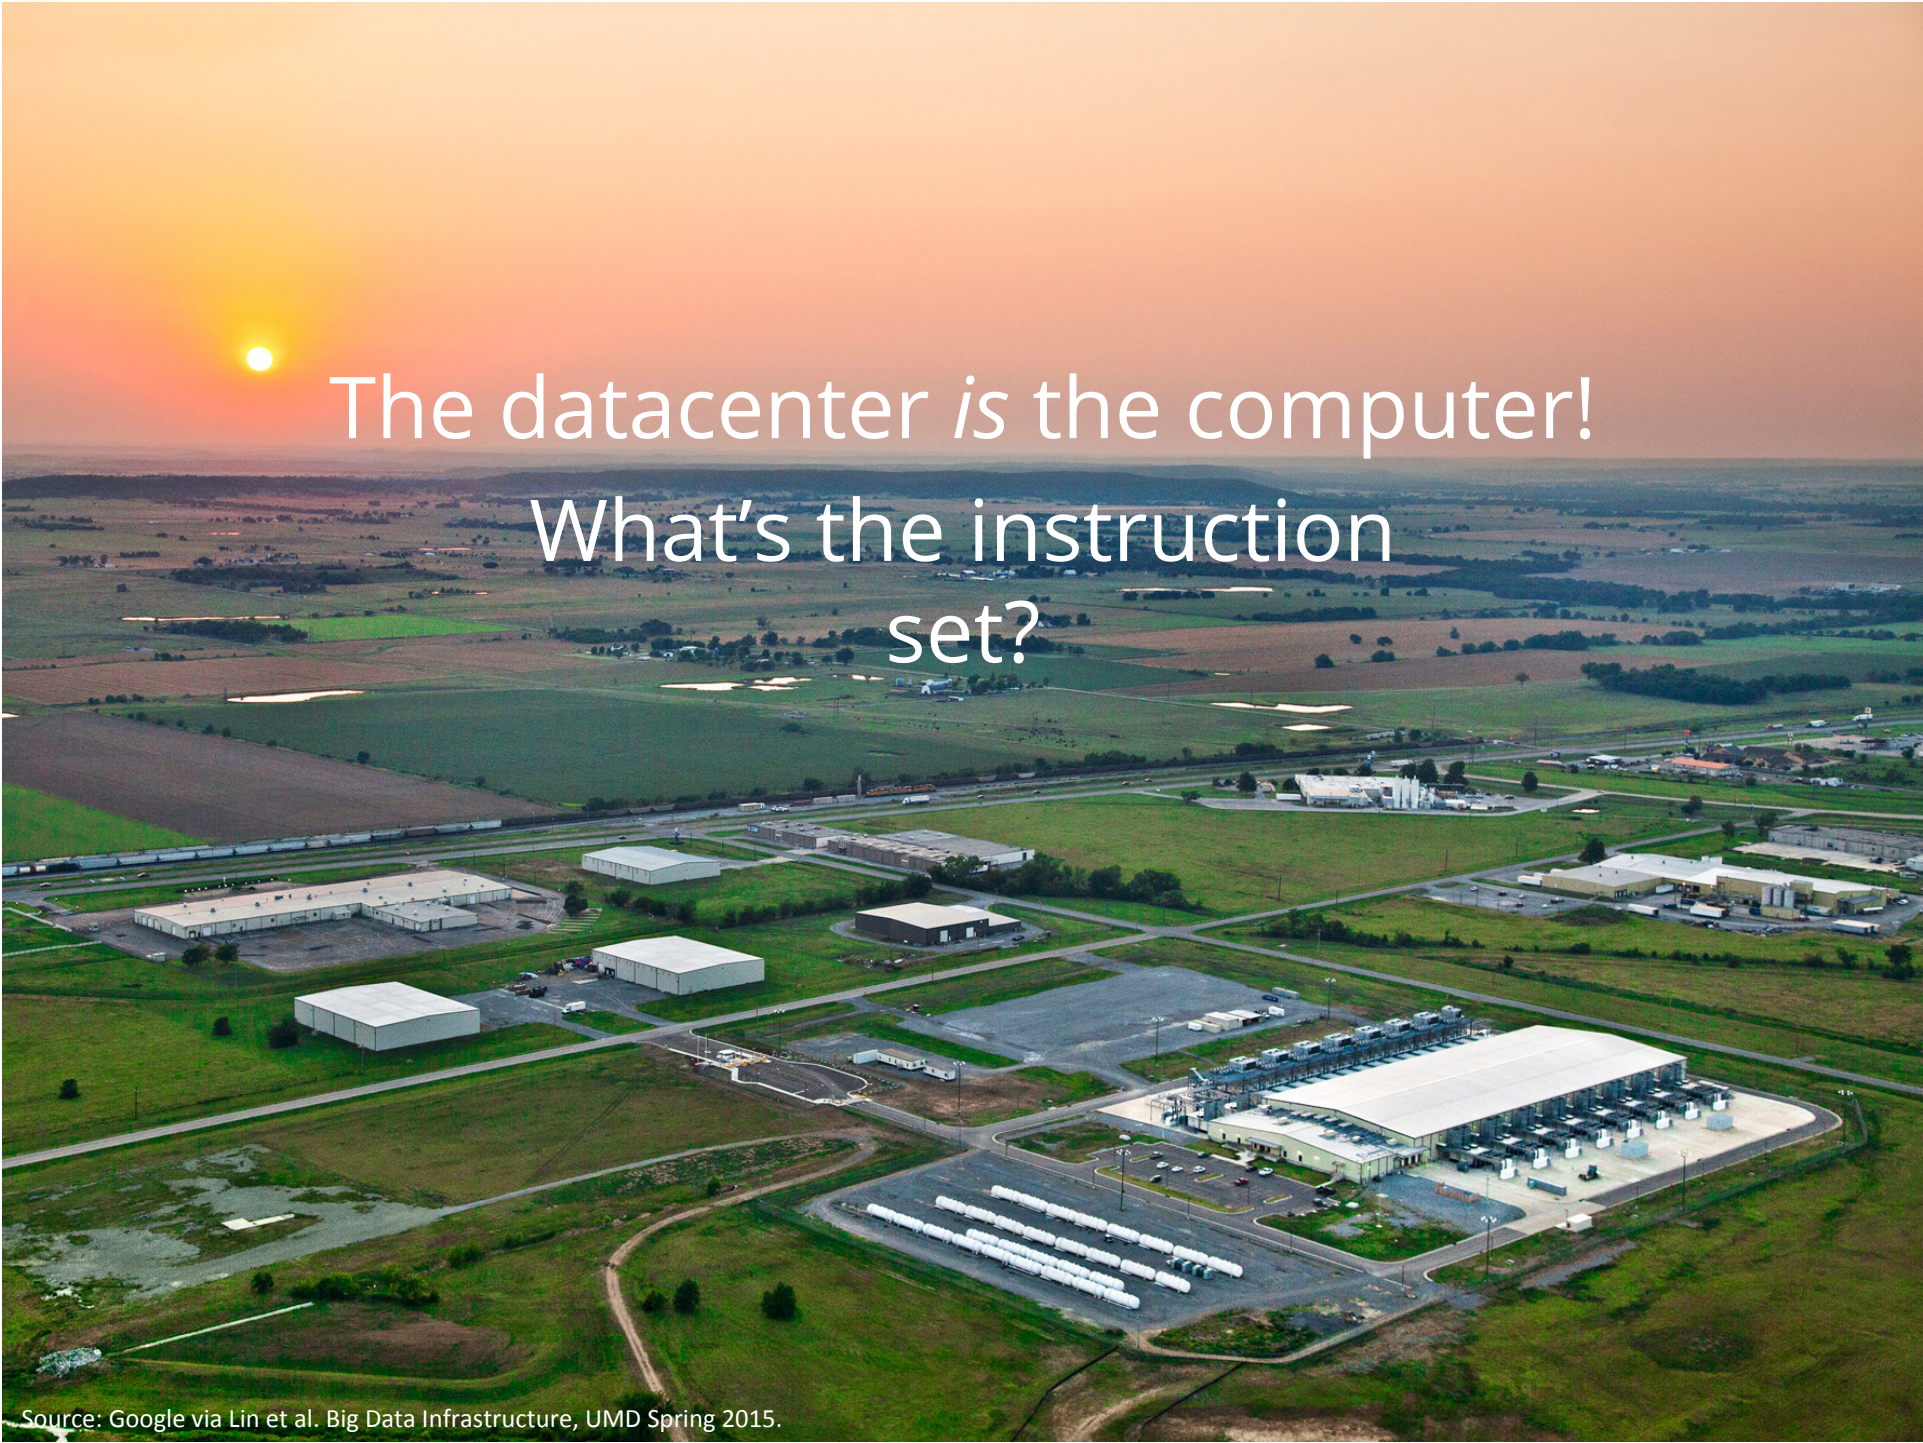


Pig Script in Hadoop



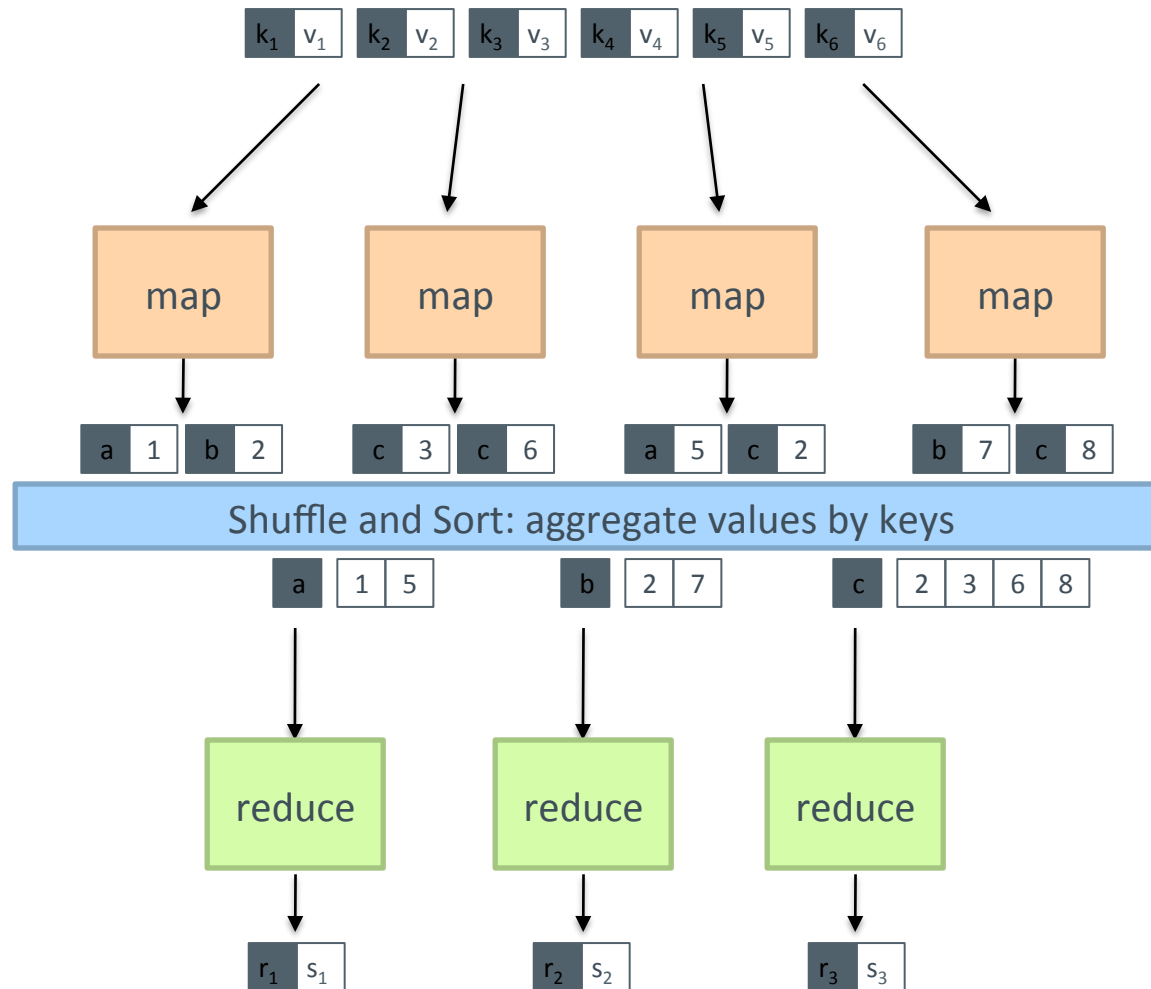


What's next?


An aerial photograph of a large industrial facility, likely a datacenter, situated in a rural area. The facility consists of several large, white, rectangular buildings with flat roofs, arranged in a grid-like pattern. In the foreground, there is a large parking lot filled with many white semi-trailers. The surrounding landscape is a mix of green fields and brown, tilled soil. In the background, there are rolling hills under a sunset sky with a bright orange and yellow glow. The sun is visible as a bright white circle on the left side of the frame.

The datacenter *is* the computer!
What's the instruction
set?

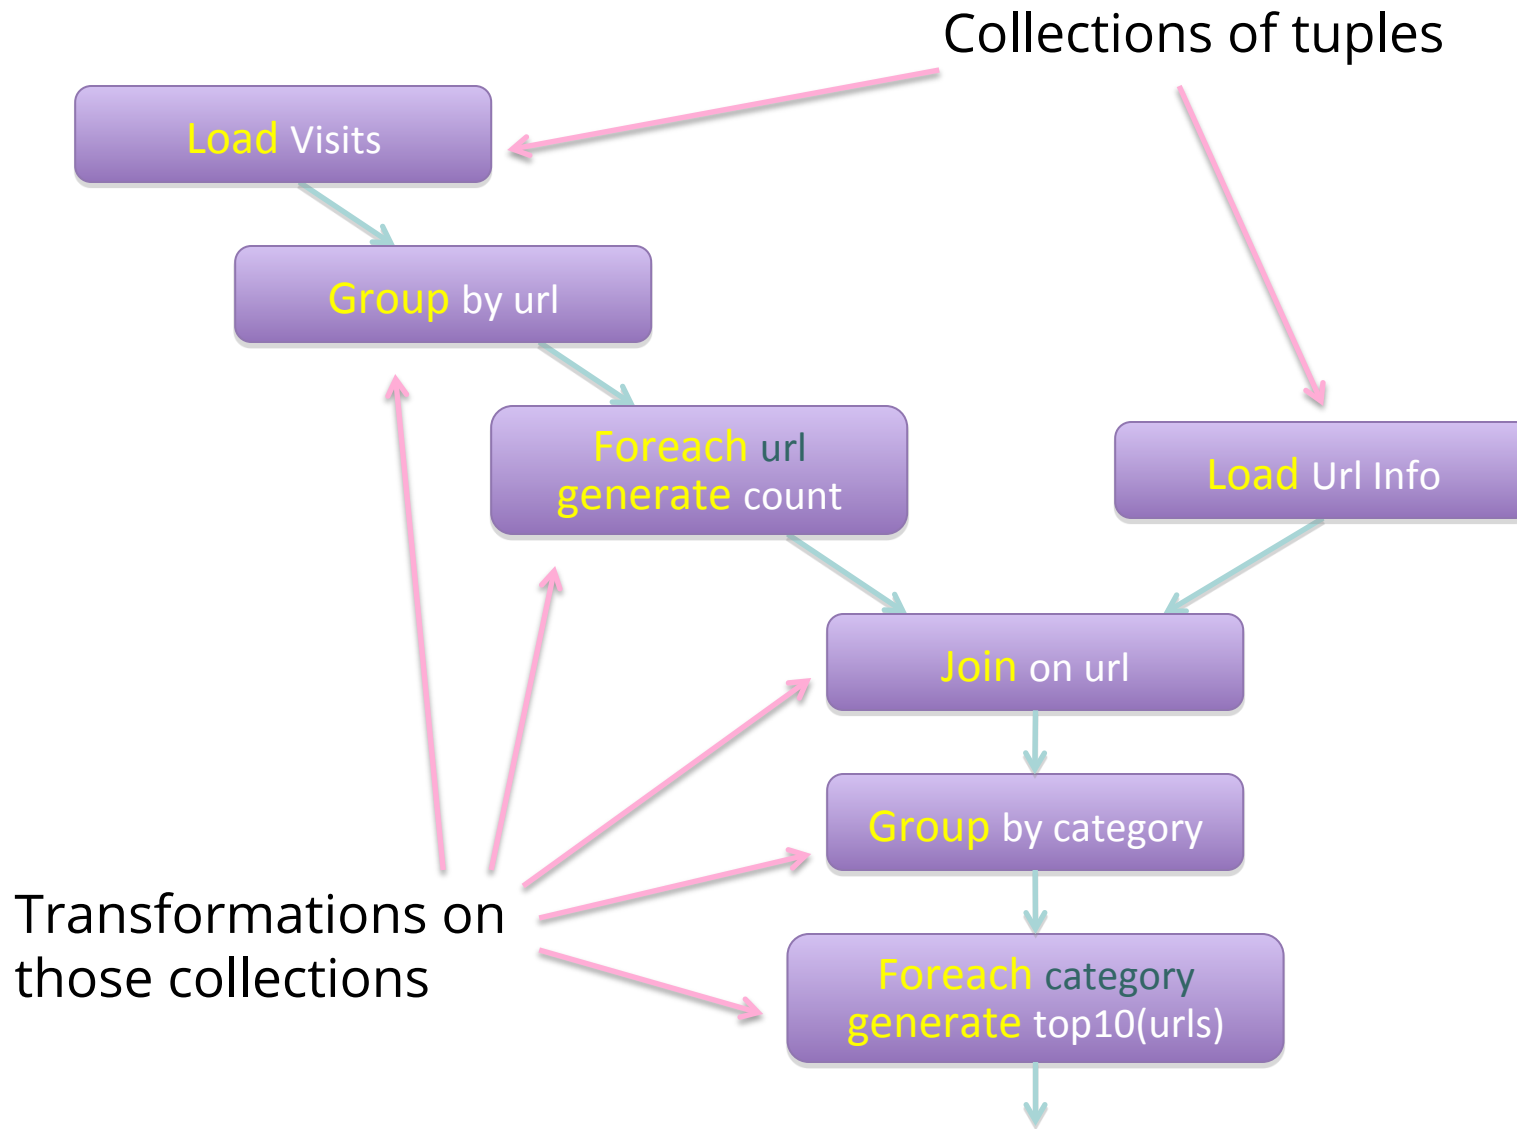
Answer?



Answer?



Generically, what is this?



Dataflows

- Comprised of:
 - Collections of records
 - Transformations on those collections
- Two important questions:
 - What are the logical operators?
 - What are the physical operators?

Spark

- One popular answer to “What’s beyond MapReduce?”
- Open-source engine for large-scale batch processing
 - Supports generalized dataflows
 - Written in Scala, with bindings in Java, Python, R
- Brief history:
 - Developed at UC Berkeley AMPLab in 2009
 - Open-sourced in 2010
 - Became top-level Apache project in February 2014
 - Commercial support provided by DataBricks

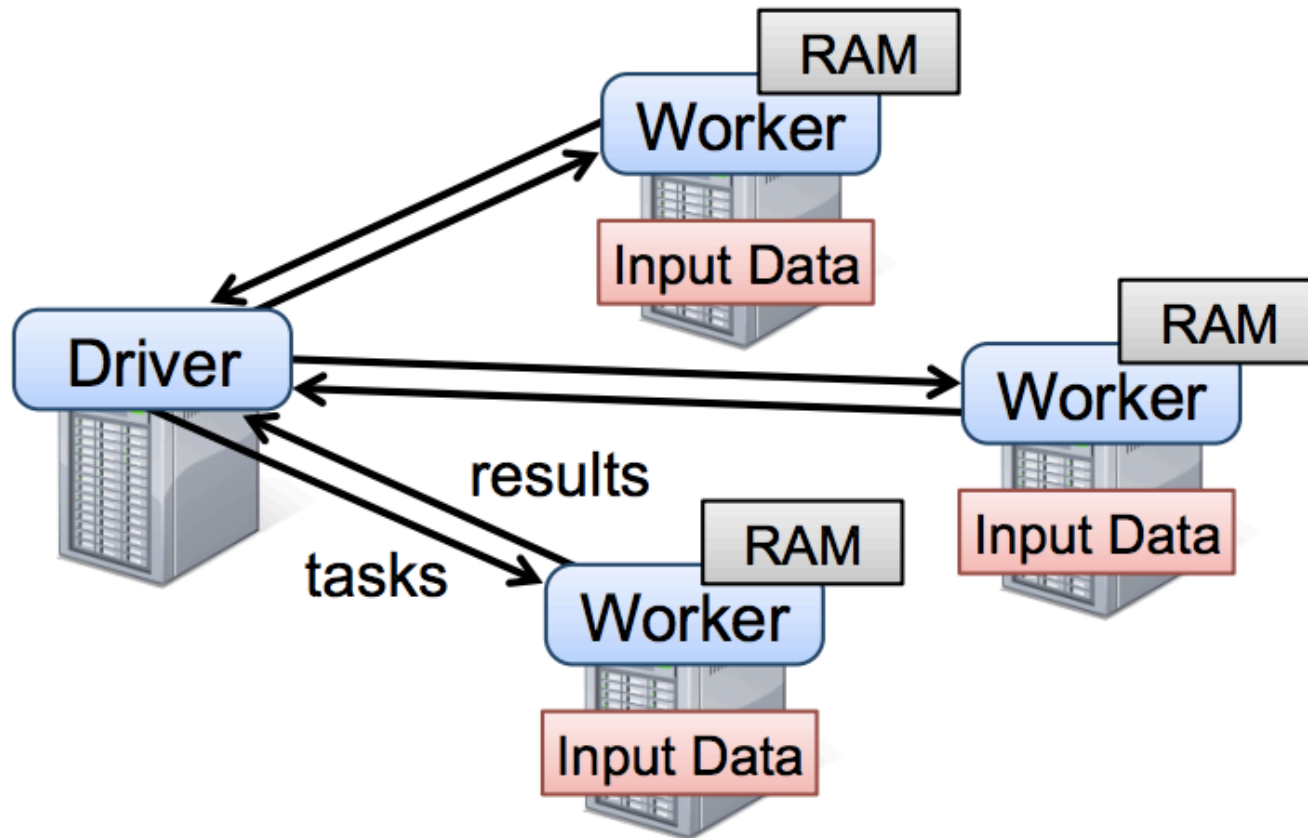
Resilient Distributed Datasets

- RDD: Spark “primitive” representing a collection of records
 - Immutable
 - Partitioned (the *D* in RDD)
- Transformations operate on an RDD to create another RDD
 - Coarse-grained manipulations only
 - RDDs keep track of *lineage*
- Persistence
 - RDDs can be materialized in memory or on disk
 - OOM or machine failures: What happens?
- Fault tolerance (the *R* in RDD):
 - RDDs can *always* be recomputed from stable storage (disk)

Operations on RDDs


- Transformations (lazy):
 - map
 - flatMap
 - filter
 - union/intersection
 - join
 - reduceByKey
 - groupByKey
 - ...
- Actions (actually trigger computations)
 - collect
 - saveAsTextFile/saveAsSequenceFile
 - ...

Spark Architecture




Spark Physical Operators


Narrow Dependencies:



map, filter




union




join with inputs
co-partitioned

Wide Dependencies:

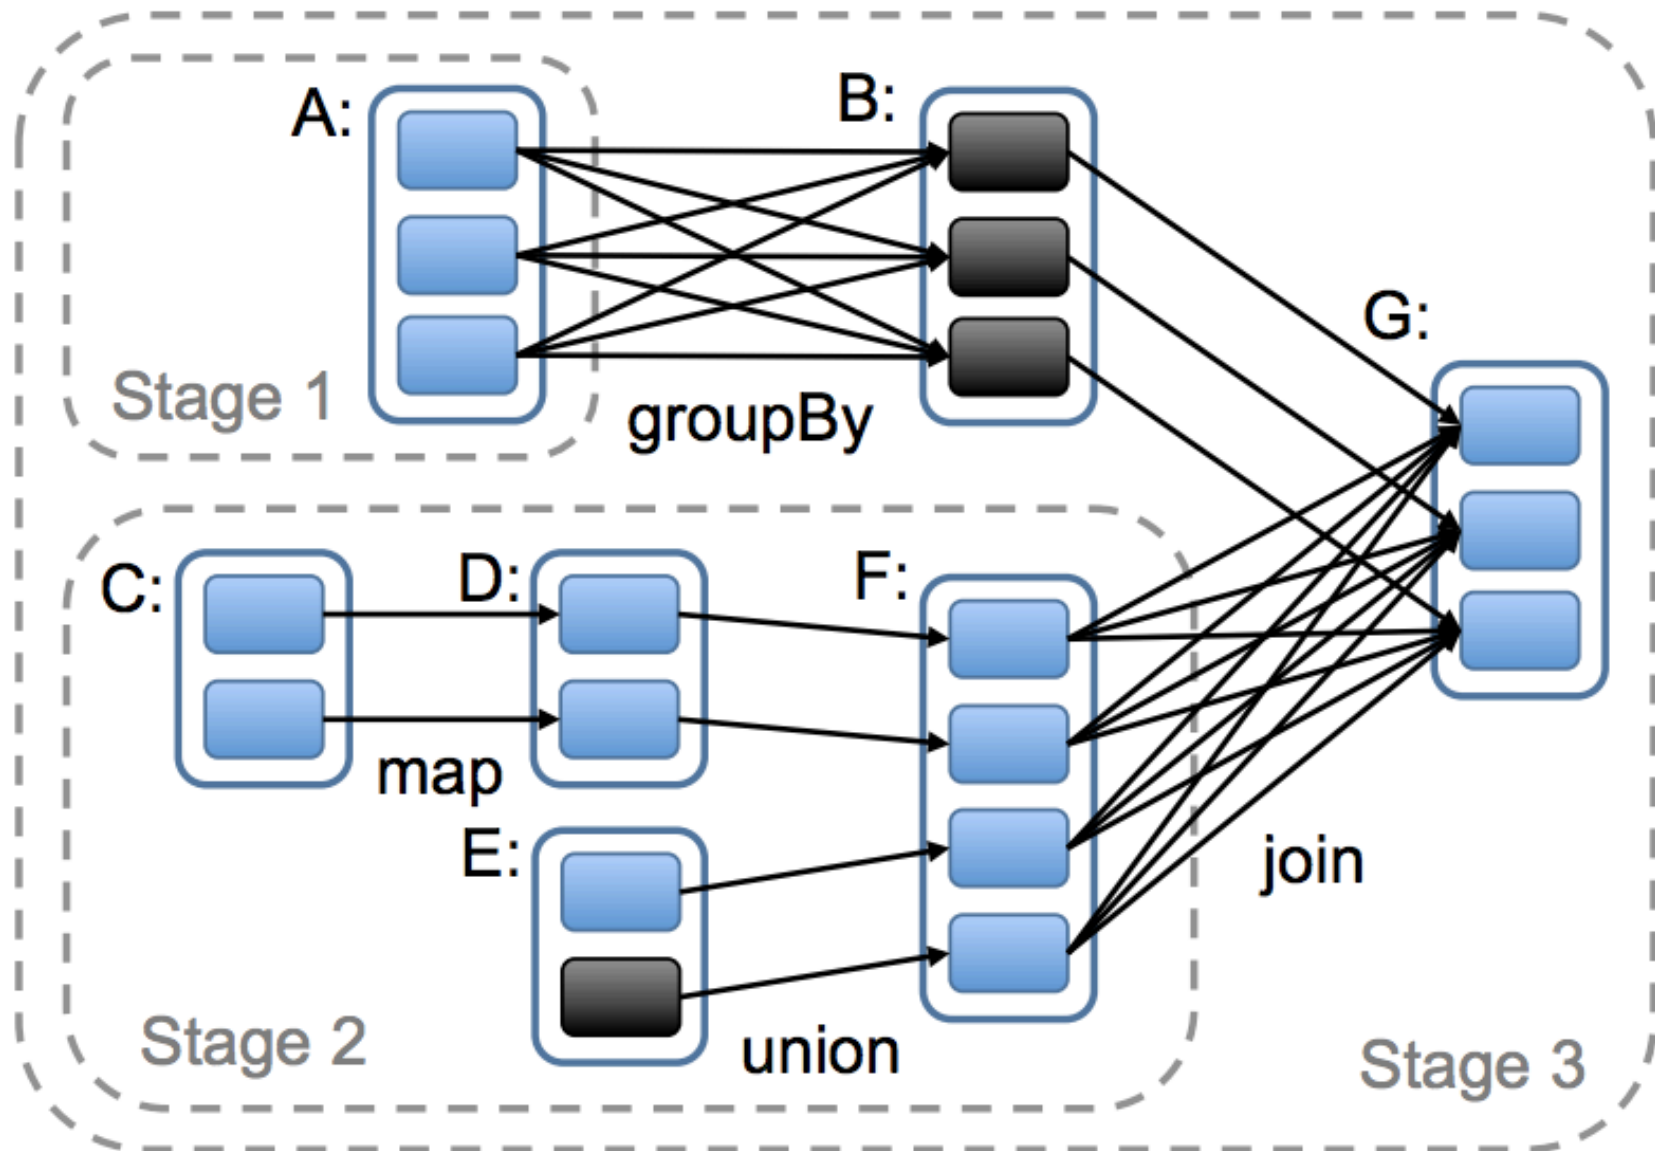


groupByKey



join with inputs not
co-partitioned

Spark Execution Plan

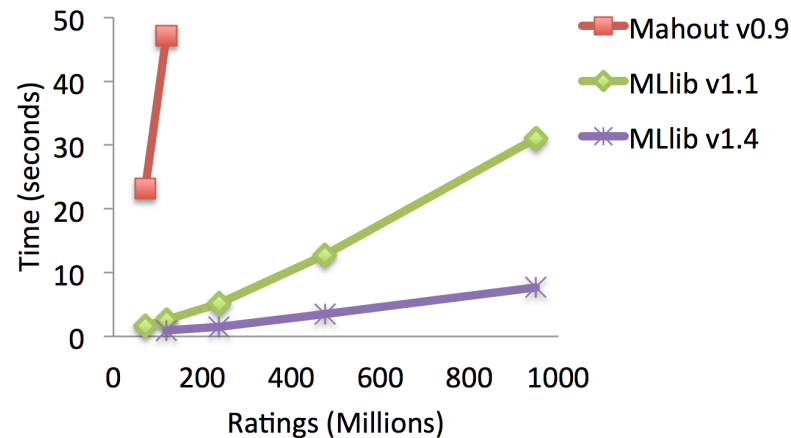


Spark DataFrames

- The hot new RDD (built on RDDs)
 - Column-oriented, schemas, ...
 - Datasets – efficient ORM
- Spark SQL
 - The hot new Shark
 - Tight integration between procedural and relational processing
 - Catalyst optimizer – “don’t bet against the compiler”
 - IndexedRDD – you can see where this is going
- GraphFrames
 - The hot new GraphX
- MLlib
 - Machine learning/fast vector math over DFs
- SparkNet ...

Spark MLlib

- Machine learning/fast vector math over dataframes
- “We observe that often a simple idea is enough: separating matrix operations from vector operations and shipping the matrix operations to be run on the cluster, while keeping vector operations local to the driver.” (Zadeh et al. 2016)



(a)

Figure 2: (a) Benchmarking results for ALS.

Questions?