

Search Engine Architecture

13. Recommender Systems

Agenda

- Recommender systems
 - Content filtering
 - Collaborative filtering
 - Nearest neighbors
 - Matrix factorization
- Semester in review

Recommender Systems

Motivation

- Contrast:
 - Hit-driven economics
 - Not enough shelf space for all CDs, DVDs
 - Not enough screens to show all movies
 - Not enough channels to show all TV programs
 - Not enough spectrum to play all music
 - Cf. online distribution
 - None of these issues!
 - We can capture the long tail of options
- From scarcity of choices to abundance...
 - A solution: recommendation engines!

Types of Recommender Systems

- Hand-curated
 - Editorial lists
- Aggregates
 - Top 10
 - Recent Uploads
- Tailored to users (another long tail)
 - Amazon
 - Pandora
 - Netflix

Two Approaches

- Content filtering – e.g., Pandora
 - Find items with content similar to other items user already likes
- Collaborative filtering – e.g., Netflix
 - Nearest neighbors
 - Find items rated highly by similar users
 - Find items rated similarly to those user already likes
 - Matrix factorization
 - Decompose ratings matrix R into PQ
 - P , Q are skinny factor loadings

Content Filtering

Content Filtering

- Create feature vector for each item
 - E.g., bag of words document-term matrix
- Create user profile vector
 - E.g., weighted average of rated items
- Score candidate items
 - E.g., cosine similarity between item and user vectors

Content Filtering

- Pros
 - No need for data on other users
 - No cold start problem for new items
 - Model is transparent – can look at features to find out why a recommendation was made
- Cons
 - Feature design requires domain expertise
 - Unable to use quality judgments from other users

Collaborative Filtering

Collaborative Filtering

- Start with ratings (a.k.a. utility) matrix:

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

Collaborative Filtering

- Nearest neighbors
 - Find items rated highly by similar users
 - Find items rated similarly to those user already likes
- Matrix factorization
 - Latent factor model
 - Decompose ratings matrix R into PQ
 - P, Q are skinny factor loadings

Collaborative Filtering: Nearest Neighbors

Nearest Neighbors

- User-user
 - Find items rated highly by similar users
 - Compute user similarity with, e.g., Pearson correlation over users' common item ratings
 - Define a user's neighborhood N of similar users
 - Then predicted rating for an item is the weighted average of ratings over user's neighborhood

Nearest Neighbors

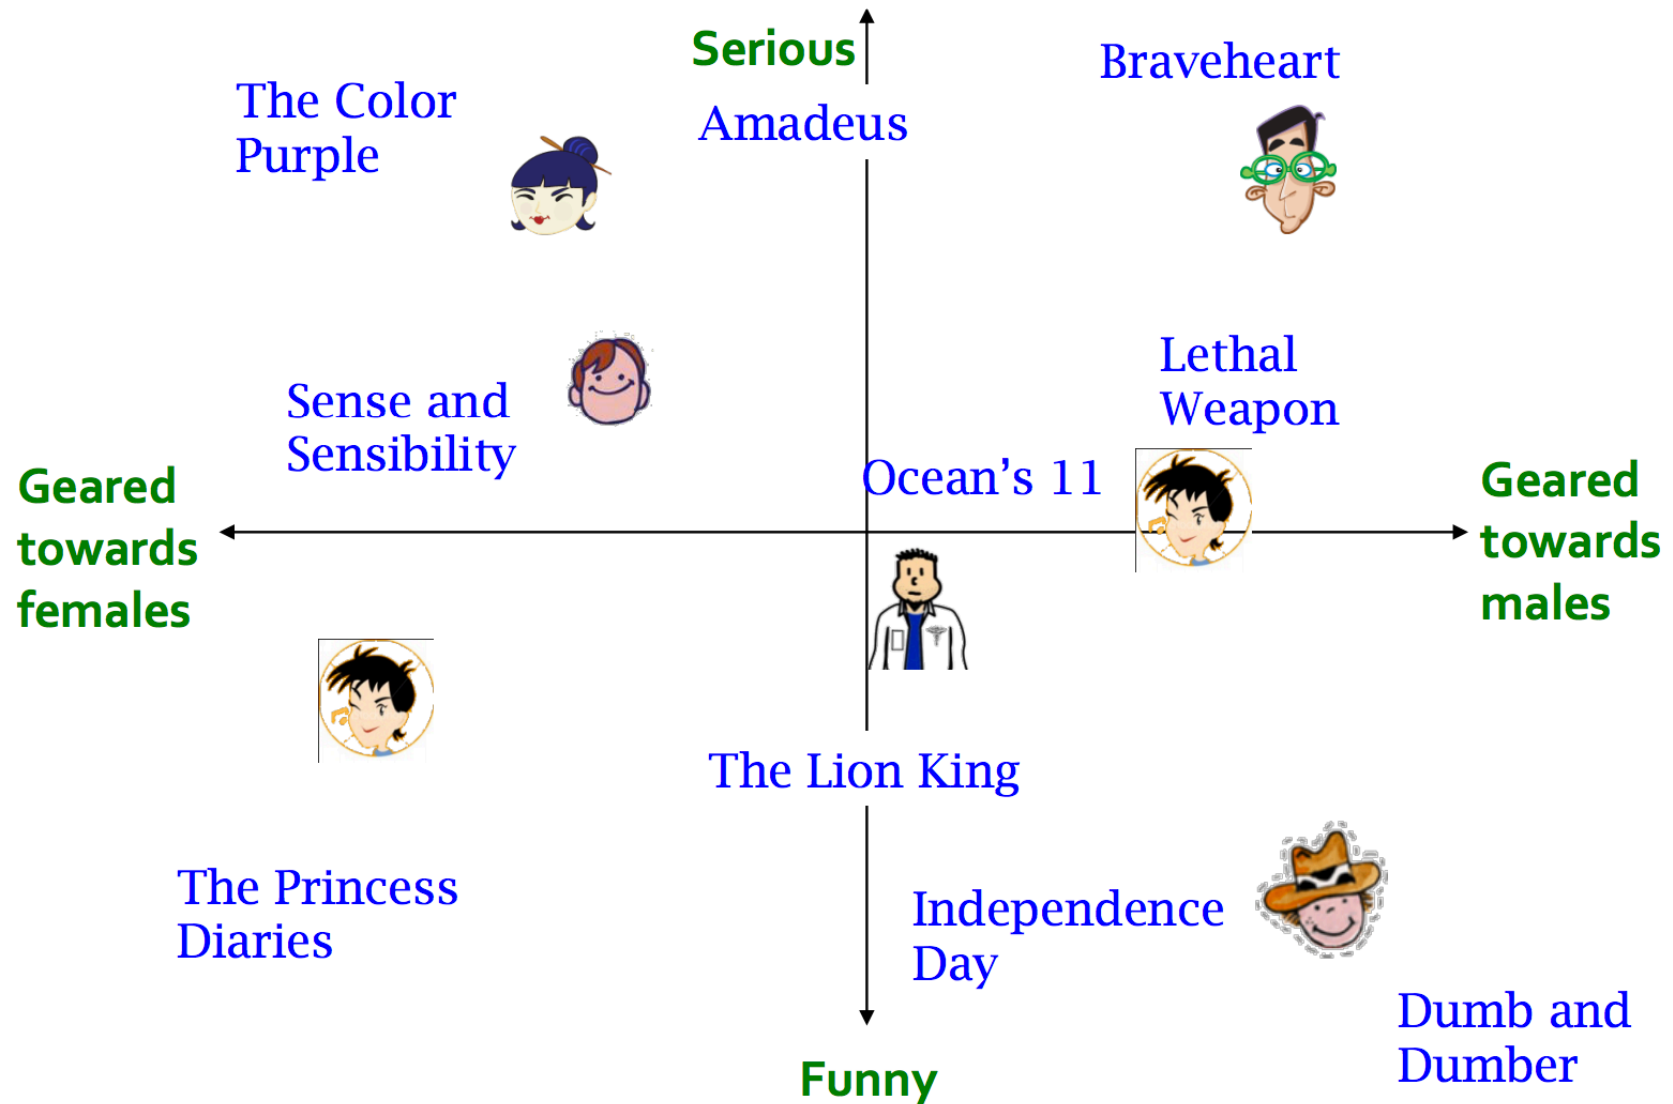
- Item-item
 - Find items similar to those rated highly
 - Compute item similarity with, e.g., Pearson correlation over common users' ratings
 - Cf. content filtering which uses item feature vector
 - Define item's neighborhood N of similar items
 - Predicted rating for an item is weighted average over item's neighborhood

Nearest Neighbors

- Pros
 - No domain expertise needed for feature design
- Cons
 - Cold start problem for new items
 - Requires users to have rated the same items
 - Problematic for sparse ratings matrix (long tail!)

Collaborative Filtering: Matrix Factorization


Latent Factor Models



SVD Recap

- Remember SVD:

- **A**: Input data matrix
- **U**: Left singular vecs
- **V**: Right singular vecs
- Σ : Singular values



- So in our case:

“SVD” on Netflix data: $R \approx Q \cdot P^T$

$$A = R, \quad Q = U, \quad P^T = \Sigma V^T$$

$$\hat{r}_{xi} = q_i \cdot p_x$$

Latent Factor Models

users

1		3		5			5		4	
		5	4			4		2	1	3
2	4		1	2		3		4	3	5
	2	4		5			4			2
		4	3	4	2				2	5
1		3		3			2			4

items

factors

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

items

users

1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

PT

factors

Q

Latent Factor Models

- Our goal is to find P and Q such that:

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2$$

The diagram illustrates the matrix factorization process. It shows a sparse matrix of observed ratings (users x items) being decomposed into a product of two matrices: a user factor matrix (users x factors) and an item factor matrix (items x factors). The observed ratings matrix is a 6x10 grid with values ranging from 1 to 5. The user factor matrix is a 6x3 grid with values ranging from -1 to 1.1. The item factor matrix is a 6x10 grid with values ranging from -1 to 2.4. The decomposition is represented by the equation: Observed Ratings = User Factors * Item Factors.

1		3			5			5		4	
		5	4			4			2	1	3
2	4		1	2		3		4	3	5	
	2	4		5			4			2	
		4	3	4	2					2	5
1		3		3		2				4	

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-.2
-1	.7	.3

1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

Overfitting

- **Want to minimize SSE for unseen test data**
- **Idea: Minimize SSE on training data**
 - Want large k (# of factors) to capture all the signals
 - But, **SSE** on test data begins to rise for $k > 2$
- This is a classical example of **overfitting**:
 - With too much freedom (too many free parameters) the model starts fitting noise
 - That is, it fits too well the training data and thus **not generalizing** well to unseen test data

1	3	4							
	3	5						5	
		4	5					5	
			3						
			3						
2				?				?	
							?		
	2	1						?	
	3								
1									

Regularization

- To solve overfitting we introduce **regularization:**

- Allow rich model where there are sufficient data
- Shrink aggressively where data are scarce

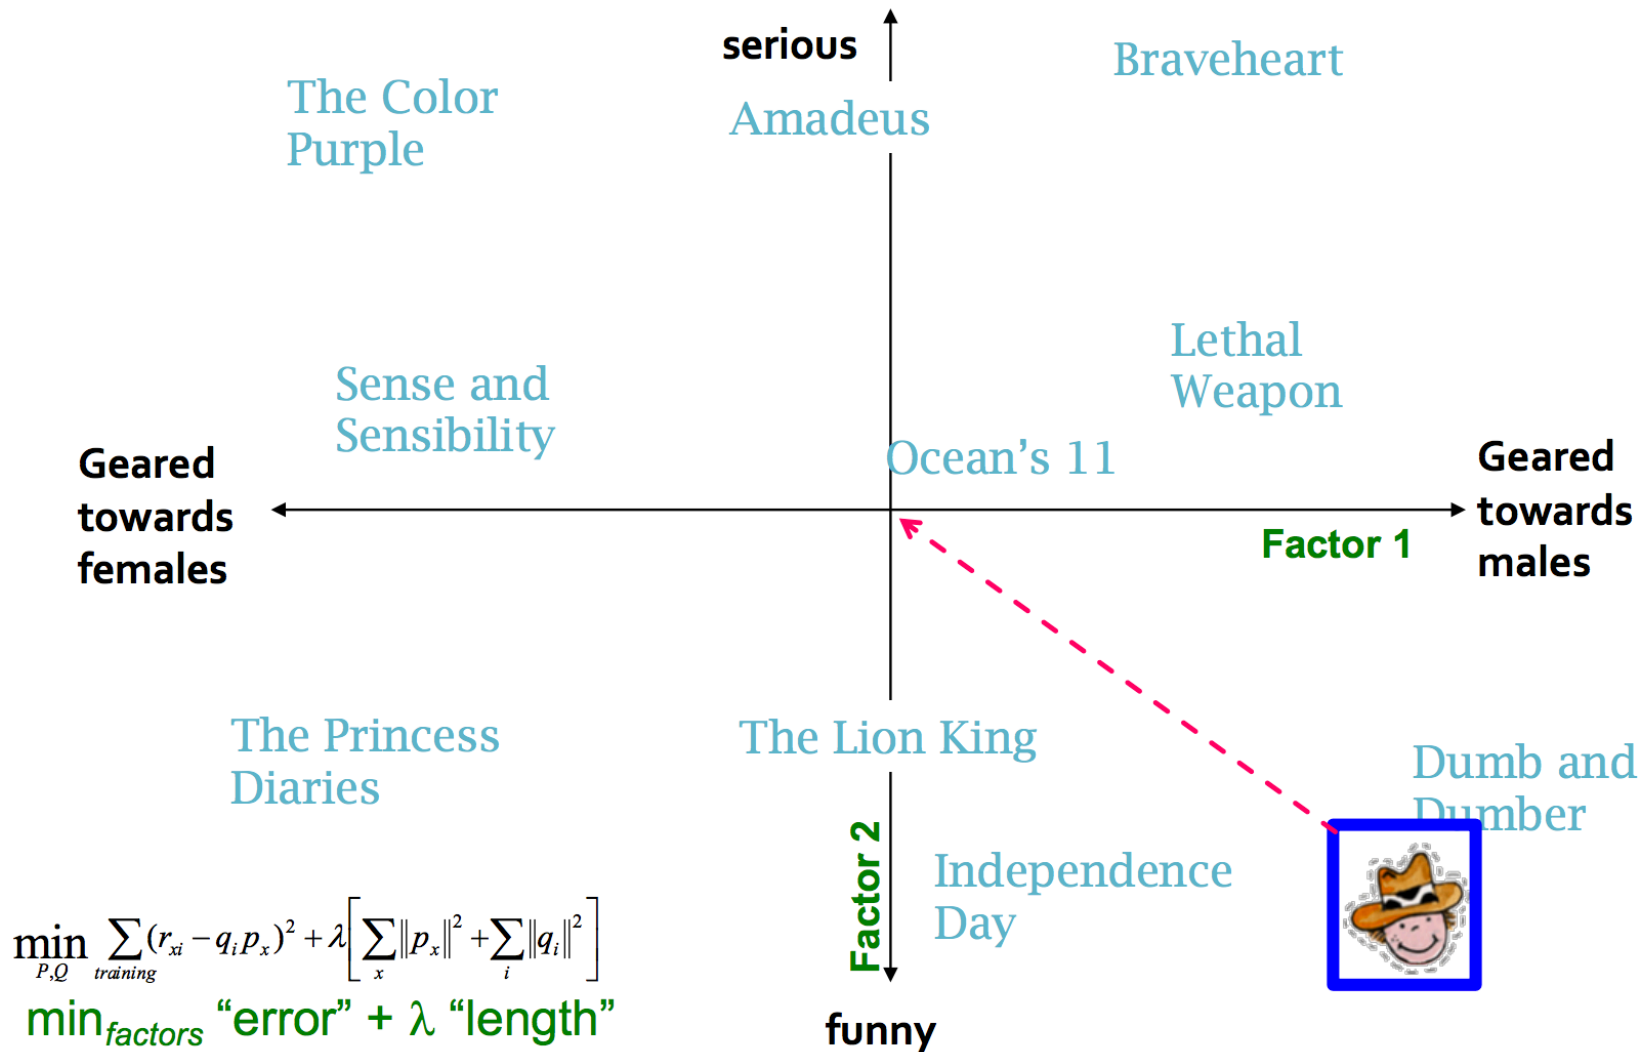
1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?	?	?
				?	?
	2	1			?
	3			?	
1					

$$\min_{P,Q} \underbrace{\sum_{training} (r_{xi} - q_i p_x)^2}_{\text{"error"}} + \underbrace{\left[\lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 \right]}_{\text{"length"}}$$

$\lambda_1, \lambda_2 \dots$ user set regularization parameters

Note: We do not care about the “raw” value of the objective function, but we care in P,Q that achieve the minimum of the objective

Effect of Regularization



Review: Recommender Systems

- Content filtering
 - Find content similar to that user already likes
- Collaborative filtering
 - Nearest neighbors
 - Find items rated highly by similar users
 - Find items rated similarly to those user already likes
 - Matrix factorization
 - Latent factor model
 - Decompose ratings matrix R into PQ
 - P, Q are skinny factor loadings
 - Item2vec? User2vec?

Semester In Review

1. Big Ideas

- Scale out, not up
- Assume failures will happen
- Good APIs hide system details
- Aim for ideal scalability
- Move code to the data
- Avoid random disk access

2. NoSQL

- Key ideas:
 - Partition – for scalability, latency
 - Replicate – for availability, throughput
 - Caching – for latency
- Key-value stores
 - Consistent hashing, hash rings
- Bigtable / LSM trees
- CAP theorem

3. Modeling and Evaluation

- Language models
- Preprocessing
 - Case folding, tokenization, stopwords, stemming
- Boolean retrieval
- Ranked retrieval
 - Vector space model, TF-IDF, cosine similarity
- Model evaluation
 - Unranked – precision, recall, F-measure
 - Ranked – MAP, nDCG

4. Indexing and Retrieval

- Inverted index
 - TF-IDF
 - Positional
- Retrieval
 - Document-at-a-time vs. term-at-a-time
- Postings list encoding (d-gaps)
- Partitioning
 - Term vs. document partitioning

5. MapReduce

- Constrained API helps with synchronization problems
- Map, combine, partition, shuffle and sort, reduce
- Data locality – pairs and stripes
- Inverted index construction
- Value-to-key conversion
- The datacenter *is* the computer!

6. Link Analysis

- Graph representation
- Shortest path
 - MapReduce – parallel BFS
- PageRank
 - Time on page under random surfer model
 - Static prior for ranking
 - Computed iteratively
- PageRank in MapReduce
 - Iterative algorithms are hard in MapReduce

7. Classification

- Supervised classification in sklearn
- Logistic regression
- Gradient descent
 - MapReduce – M partial gradients, 1 model update
- Stochastic gradient descent
- Ensemble methods
 - MapReduce implementation – mappers only
- Case study: GoogLeNet 2014

8. Clustering

- For exploratory analysis, recommender systems, preprocessing, ...
- Hierarchical agglomerative clustering
 - Start with N clusters, merge until one
- K-means
 - Iteratively recompute centroids and reassign points
 - MapReduce – map: assign, reduce: new centroids
- Gaussian mixture models
 - Soft assignment of points to clusters
 - MapReduce – similar to K-means

9. Distributed Word Representations

- Distributed representations
/ distributional hypothesis
- Dimensionality reduction
- Artificial neural networks
- Representation learning
- Word2vec
 - Skip-gram
 - CBOW
- Doc2vec
- SVD reduction

10. Learning to Rank

- ML vs. IR
- Classification
 - Predict class of query-document pair
- Pointwise learning
 - Learn thresholds to separate ranks
- Pairwise learning
 - Turns ordinal regression into binary classification
- Issues
 - Cost sensitivity for high-ranked documents
 - Query normalization

11. Beyond MapReduce

- Addressing MapReduce criticisms
 - Schemas with Thrift
 - High-level languages – Hive, Pig
- Dataflow – DAG of transformations
- Spark
 - RDD – store transforms needed to reproduce data
- Pregel
 - Graph-centric, express graph algorithms naturally
 - Each vertex passes messages to neighbors
 - Synchronization via supersteps

12. Streams

- Sampling
- Hashing
 - Set cardinality – HyperLogLog counter
 - Set membership – Bloom filter
 - Frequency estimation – Count-min sketch
- Storm
 - Spouts, bolts, and clever tracking
- Spark Streaming
 - Small, deterministic batch jobs
- Dataflow
 - Windows, triggers, and incremental processing

13. Finding Similar Items

- Represent documents with short signatures
 - Minhash
 - Given hash function, find term with smallest hash value
 - $P[h_1(D_1) = h_2(D_2)] = \text{Jaccard}(D_1, D_2)$
- Find candidates that are likely similar
 - Compute k minhashes per document (“band”)
 - Documents that match in a band are candidates
 - Evaluate candidates thoroughly
 - Repeat for n bands

14. Recommender Systems

- Content filtering
 - Find content similar to that user already likes
- Collaborative filtering
 - Nearest neighbors
 - Find items rated highly by similar users
 - Find items rated similarly to those user already likes
 - Matrix factorization
 - Latent factor model
 - Decompose ratings matrix R into PQ
 - P, Q are skinny factor loadings

Thank you!